

# What is data engineering

Data engineering is the development, implementation, and maintenance of systems and processes that take in raw data and produce high-quality, consistent information that supports downstream use cases, such as analysis and machine learning. Data engineering is the intersection of security, data management, DataOps, data architecture, orchestration, and software engineering. A data engineer manages the data engineering lifecycle, beginning with getting data from source systems and ending with serving data for use cases, such as analysis or machine learning.

*-Fundamentals of Data Engineering*

# What is a relational database?

- a relation (table) is a collection of tuples. Each tuple is called a \*row\*
- a database is a collection of tables related to each other through common data values.
- Everything in a column is values of one attribute
- A cell is expected to be atomic, no lists, dictionaries, etc
- Tables are related to each other if they have columns called keys which represent the same values
- SQL a declarative model: a query optimizer decides how to execute the query (if a field range covers 80% of values, should we use the index or the table?). Also parallelizable

# How would you model data?

- The needs of OLTP databases are very different from those of OLAP databases
- OLTP databases usually need CRUD operations: CReate, Update, Delete
- OLTP tables (and incoming OLAP schemas) have a star like structure. *Fact tables* with pointers, or **keys** to *dimension tables*.
- Normalization: *The attributes of a table should be dependent on the primary key, on the whole key and nothing but the key.*

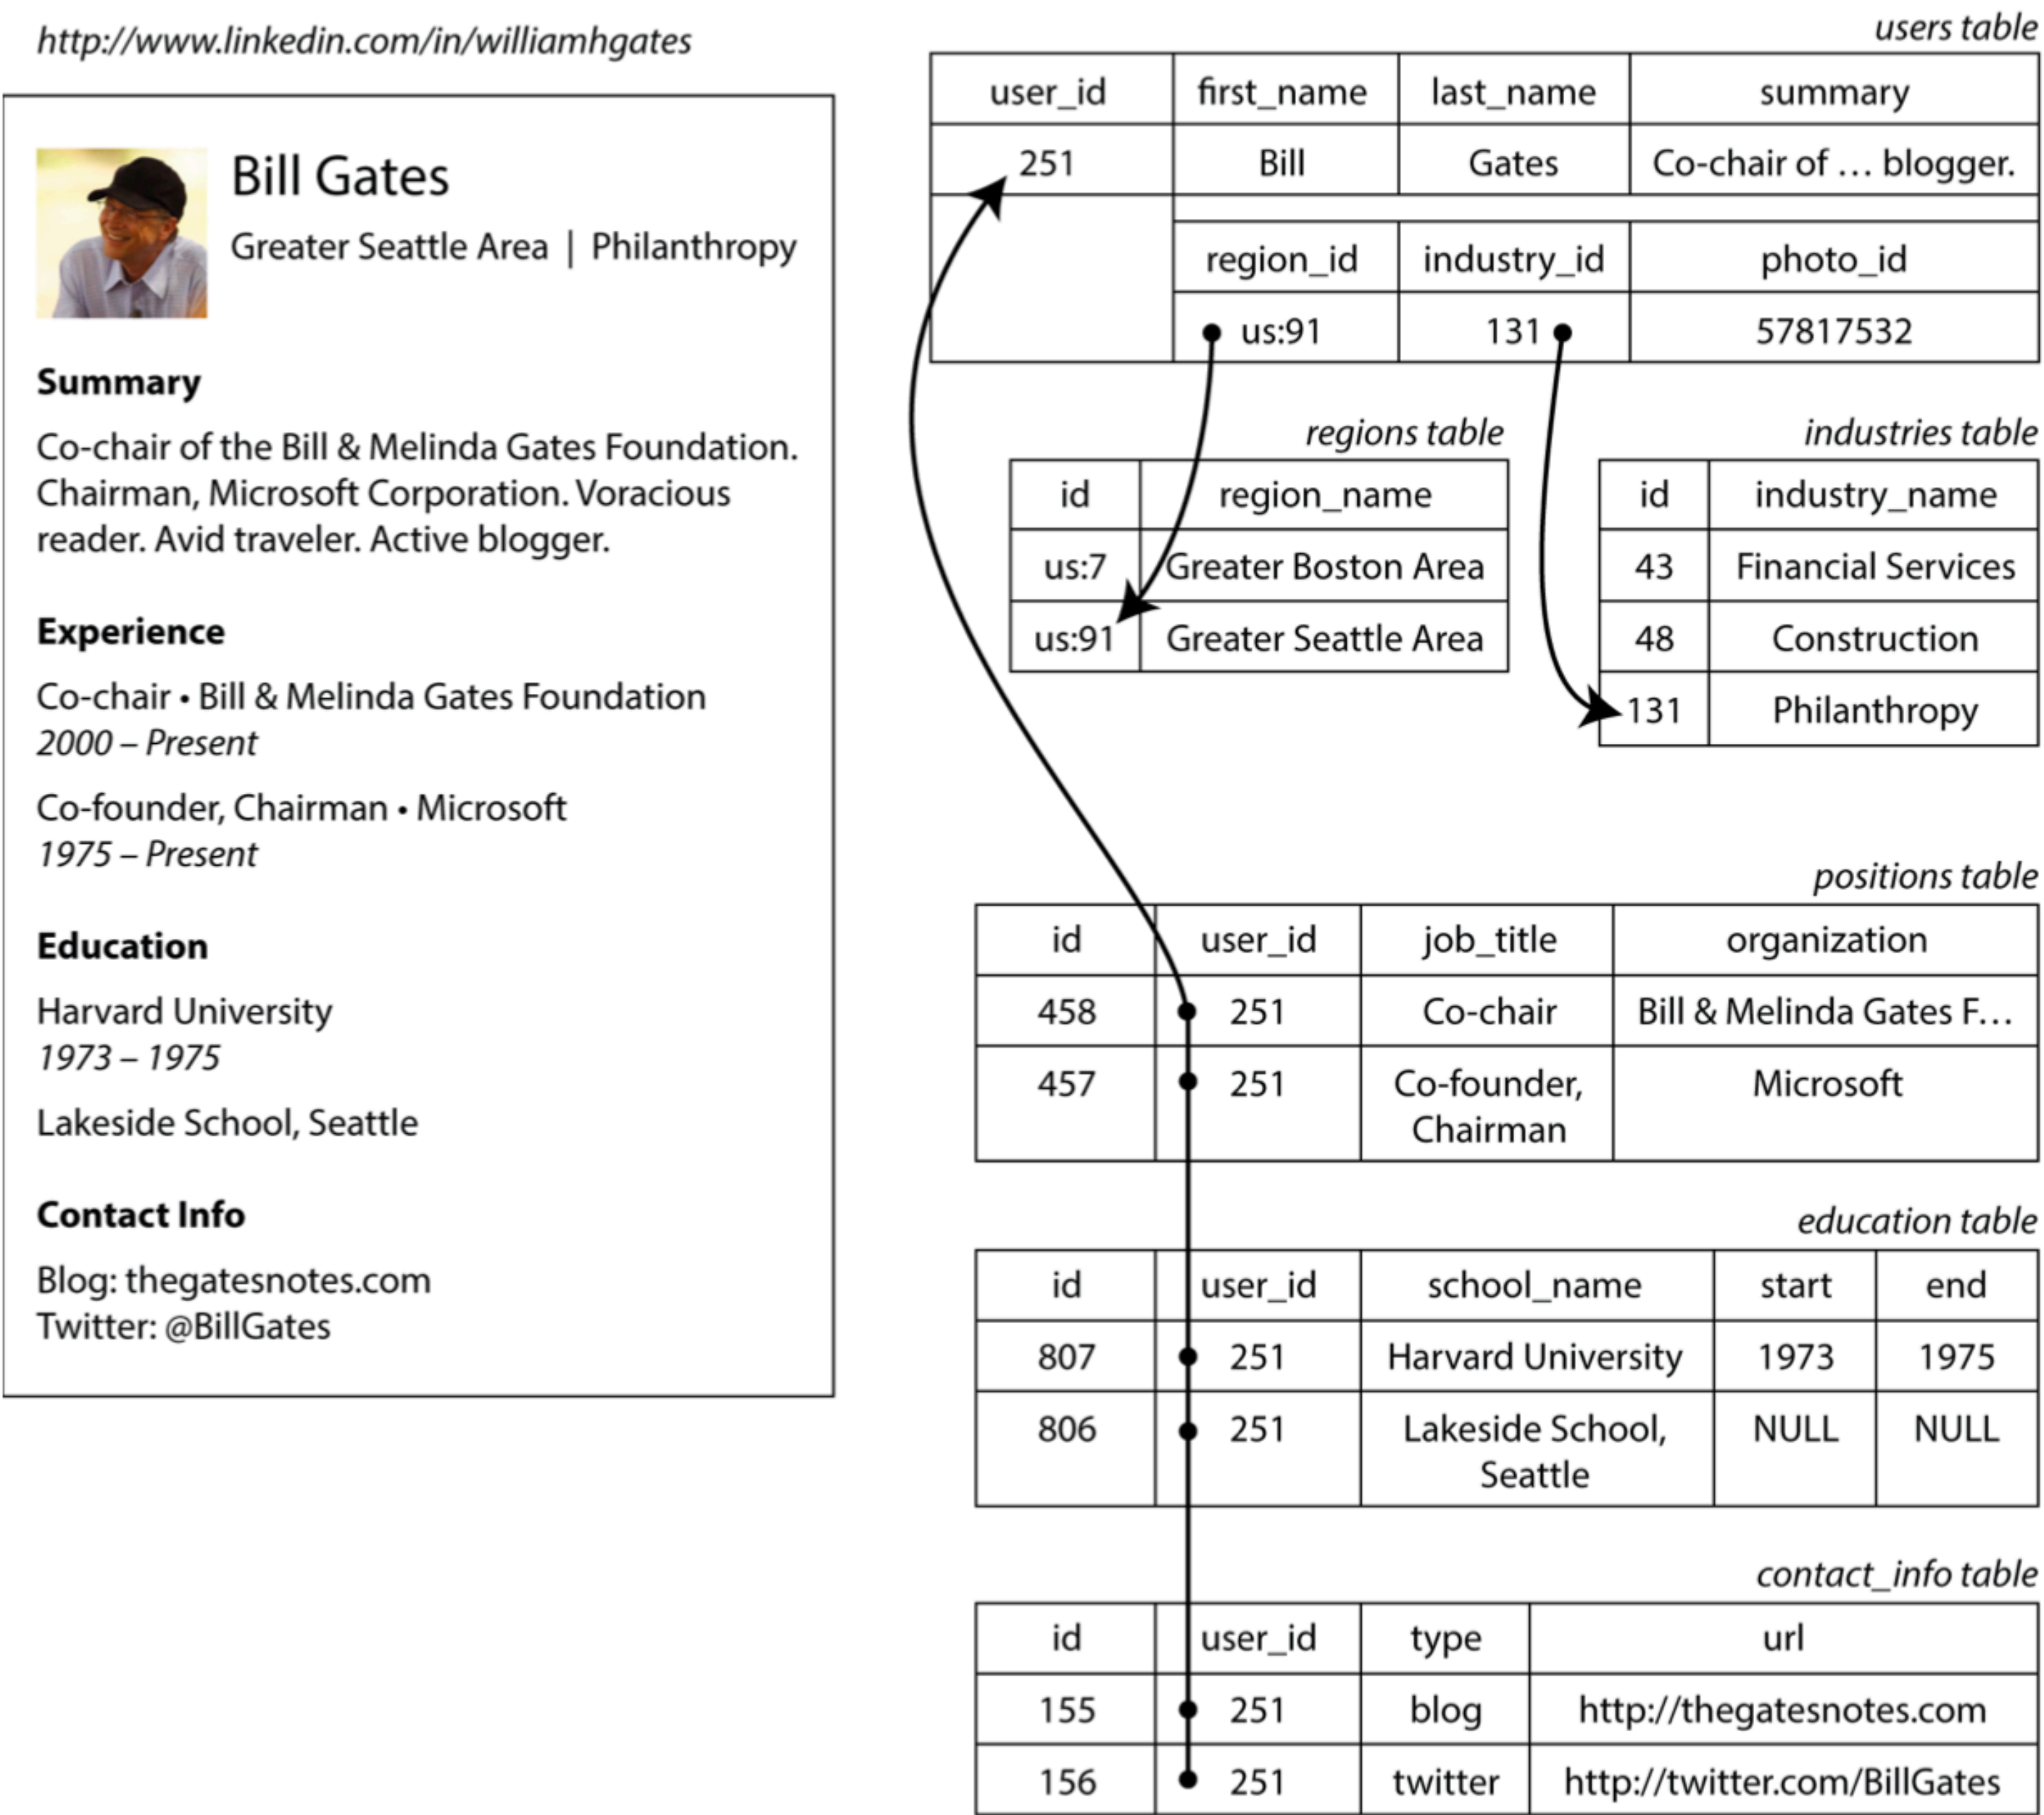





Table: New RecordDelete Record


	id	last_name	first_name	middle_name	street_1	street_2	city	state	zip	amount	date	candidate_id
	Filter	Filter	Filter	Filter	Filter	Filter	Filter	Filter	Filter	Filter	Filter	Filter
1	1	Agee	Steven	NULL	549 Laurel ...	NULL	Floyd	VA	24091	500	2007-06-30	16
2	5	Akin	Charles	NULL	10187 Suga...	NULL	Bentonville	AR	72712	100	2007-06-16	16
3	6	Akin	Mike	NULL	181 Baywo...	NULL	Monticello	AR	71655	1500	2007-05-18	16
4	7	Akin	Rebecca	NULL	181 Baywo...	NULL	Monticello	AR	71655	500	2007-05-18	16
5	8	Aldridge	Brittni	NULL	808 Capitol...	NULL	Washington	DC	20024	250	2007-06-06	16
6	9	Allen	John D.	NULL	1052 Cann...	NULL						
7	10	Allen	John D.	NULL	1052 Cann...	NULL						
8	11	Allison	John W.	NULL	P.O. Box 10...	NULL						
9	12	Allison	Rebecca	NULL	3206 Sum...	NULL						

	id	first_name	last_name	middle_name	party
	Filter	Filter	Filter	Filter	Filter
1	16	Mike	Huckabee		R
2	20	Barack	Obama		D
3	22	Rudolph	Giuliani		R
4	24	Mike	Gravel		D
5	26	John	Edwards		D
6	29	Bill	Richardson		D
7	30	Duncan	Hunter		R
8	31	Dennis	Kucinich		D
9	32	Ron	Paul		R

# Create Tables

```
  
  
DROP TABLE IF EXISTS "candidates";  
DROP TABLE IF EXISTS "contributors";  
CREATE TABLE "candidates" (  
    "id" INTEGER PRIMARY KEY NOT NULL ,  
    "first_name" VARCHAR,  
    "last_name" VARCHAR,  
    "middle_name" VARCHAR,  
    "party" VARCHAR NOT NULL  
);  
CREATE TABLE "contributors" (  
    "id" INTEGER PRIMARY KEY AUTOINCREMENT NOT NULL,  
    "last_name" VARCHAR,  
    "first_name" VARCHAR,  
    "middle_name" VARCHAR,  
    "street_1" VARCHAR,  
    "street_2" VARCHAR,  
    "city" VARCHAR,  
    "state" VARCHAR,  
    "zip" VARCHAR, -- Notice that we are converting the zip from integer to string  
    "amount" INTEGER,  
    "date" DATETIME,  
    "candidate_id" INTEGER NOT NULL,  
    FOREIGN KEY(candidate_id) REFERENCES candidates(id)  
);
```

# selects in tables



```
SELECT * FROM contributors WHERE amount BETWEEN 20 AND 40;
SELECT * FROM contributors WHERE state='VA' AND amount < 400;
SELECT * FROM contributors WHERE state IN ('VA', 'WA');
SELECT * FROM contributors WHERE state IS NULL;
SELECT * FROM contributors WHERE state IS NOT NULL;
SELECT * FROM contributors ORDER BY amount;
SELECT * FROM contributors ORDER BY amount DESC;
SELECT * FROM contributors ORDER BY amount DESC LIMIT 10;
SELECT AVG(amount) FROM contributors;
SELECT state,AVG(amount) FROM contributors GROUP BY state;
```



# inserts and alters


```
INSERT INTO candidates (id, first_name, last_name, middle_name, party) VALUES (?, ?, ?, ?, ?);
ALTER TABLE contributors ADD COLUMN name;
ALTER TABLE contributors DROP COLUMN name;
DELETE FROM contributors WHERE last_name="Ahrens";
drop table if exists mailing_list;

create table mailing_list (
    email          varchar(100) not null primary key,
    name           varchar(100)
);

drop table if exists phone_numbers;
create table phone_numbers (
    email          varchar(100) not null references mailing_list(email),
    number_type    varchar(15) check (number_type in ('work', 'home', 'cell', 'beeper')),
    phone_number   varchar(20) not null
);

insert into phone_numbers values ('ogrady@fastbuck.com', 'work', '(800) 555-1212');
insert into phone_numbers values ('ogrady@fastbuck.com', 'home', '(617) 495-6000');
insert into phone_numbers values ('philg@mit.edu', 'work', '(617) 253-8574');
insert into phone_numbers values ('ogrady@fastbuck.com', 'beeper', '(617) 222-3456');
```

# Inner Joins



left					right					Result						
	key1	key2	A	B		key1	key2	C	D		key1	key2	A	B	C	D
0	K0	K0	A0	B0	0	K0	K0	C0	D0	0	K0	K0	A0	B0	C0	D0
1	K0	K1	A1	B1	1	K1	K0	C1	D1	1	K1	K0	A2	B2	C1	D1
2	K1	K0	A2	B2	2	K1	K0	C2	D2	2	K1	K0	A2	B2	C2	D2
3	K2	K1	A3	B3	3	K2	K0	C3	D3							





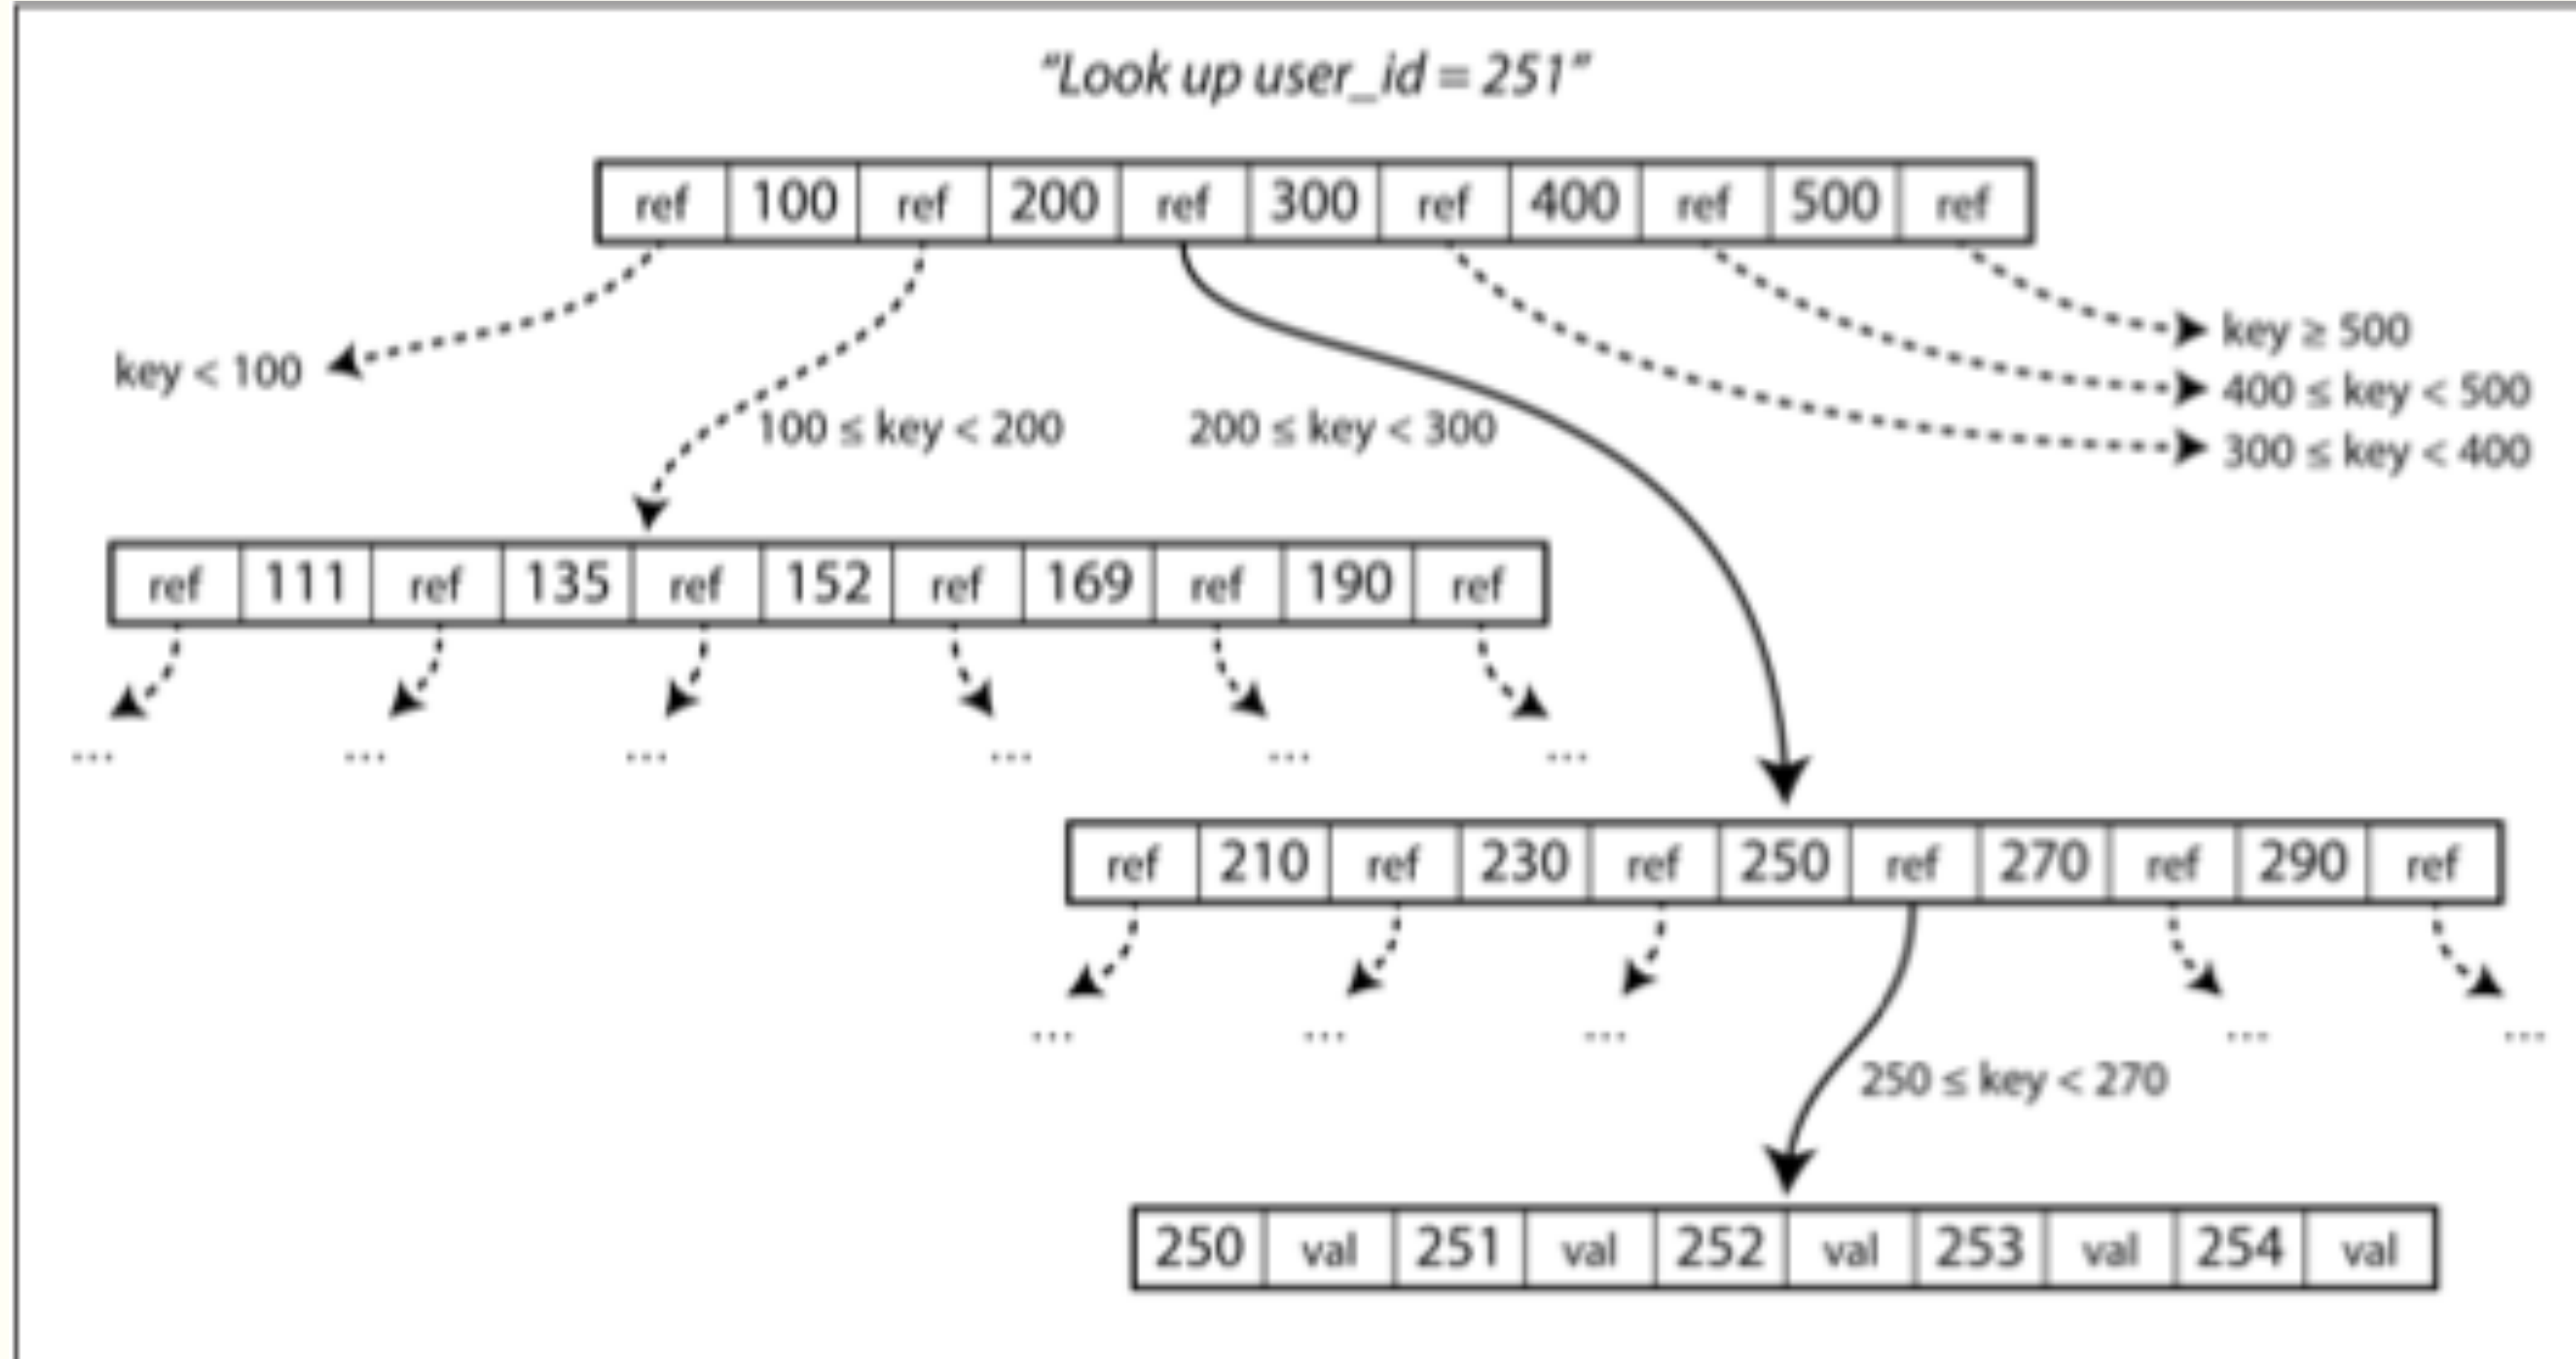
# Storage Components of a RDBMS

- the heap file: this is where the rows or columns are stored
- regular relational databases use row oriented heap files
- an index file(s): this is where the index for a particular attribute is stored
- sometimes you have a clustered index ( all data stored in index) or covering index (some data is stored in index)
- the WAL or write-ahead log: this is used to handle transactions

# Binary Search Trees

- These have the property that for any value  $x$ , numbers less go to the left and numbers right go to the right
- Consider a list: [17,5,35,2,11,29,38,9,16,7,8]
- Then the binary tree you get is:





# Btrees



### Meta Page

Pgno: 0  
Misc...  
Root : 1

### Data Page

Pgno: 1  
Misc...  
offset: 4000  
offset: 3000  
2,bar  
1,foo

### Write-Ahead Log

Add 1,foo to  
page 1  
Commit  
Add 2,bar to  
page 1

### Meta Page

Pgno: 0  
Misc...  
Root : 1

### Data Page

Pgno: 1  
Misc...  
offset: 4000  
offset: 3000  
2,bar  
1,foo

### Write-Ahead Log

Add 1,foo to  
page 1  
Commit  
Add 2,bar to  
page 1  
Commit

### Meta Page

Pgno: 0  
Misc...  
Root : 1

### Data Page

Pgno: 1  
Misc...  
offset: 4000  
offset: 3000  
2,bar  
1,foo

### Write-Ahead Log

Add 1,foo to  
page 1  
Commit  
Add 2,bar to  
page 1  
Commit  
Checkpoint

### Meta Page

Pgno: 0  
Misc...  
Root : 1

### Data Page

Pgno: 1  
Misc...  
offset: 4000  
offset: 3000  
2,bar  
1,foo

# WAL

# General requirements for e-commerce

- The batch of operations is viewed as a single atomic operation, so all of the operations either succeed together or fail together.
- The database is in a valid state before and after the transaction.
- The batch update appears to be isolated; no other query should ever see a database state in which only some of the operations have been applied.

# Transactions: ACID

- A is for **atomicity**. The batch of operations is viewed as a single atomic operation, so all of the operations either succeed together or fail together. This means that the batch of operations either all happen (\*\*commit\*\*) or not happen at all (\*\*abort\*\*, \*\*rollback\*\*).
- The batch update appears to be **isolated**; other queries should never see a database state in which only some of the operations have been applied. I is for Isolation. This *is the most interesting of the lot*, and critical to the sensible running of a database. The idea is that transactions should not step on each other. Each transaction should pretend that it's the only one running on the database: in other words, as if the transactions were completely serialized.



- In practice this would make things very slow, so we try different transactional guarantees that fall short of explicit serialization except in the situations that really need serialization.
- The database is in a valid state before and after the transaction. D is for **Durability**: once a transaction has committed successfully, data committed won't be forgotten. This requires persistent storage, or replication, or both.
- C is for **Consistency**: data invariants must be true. This is really a property of the application: eg accounting tables must be balanced. Databases can help with foreign keys, but this is a property of the app. We won't discuss this one further.


# Why is isolation important?

- It is hard to program without isolation. isolation is what guarantees stability. It is what makes sure that there are no dirty reads and dirty writes.
- **Dirty reads:** One client reads another client's writes before they have been committed. This could mean you see a value that would be later rolled back.
- **Dirty writes:** One client overwrites data that another client has written, but not yet committed. Bad. When we write we will use a lock to ensure no-one else can write.
- Clearly, the notions of isolation are really the notions of concurrency: these issues will also occur when 2 programs access any data, in memory or in a database. In both cases locks and other ideas must be used to make sure that there is only one mutator at a time, and that an object is not exposed in an inconsistent state.


**From transactions to analysis**



# Organizational Evolution

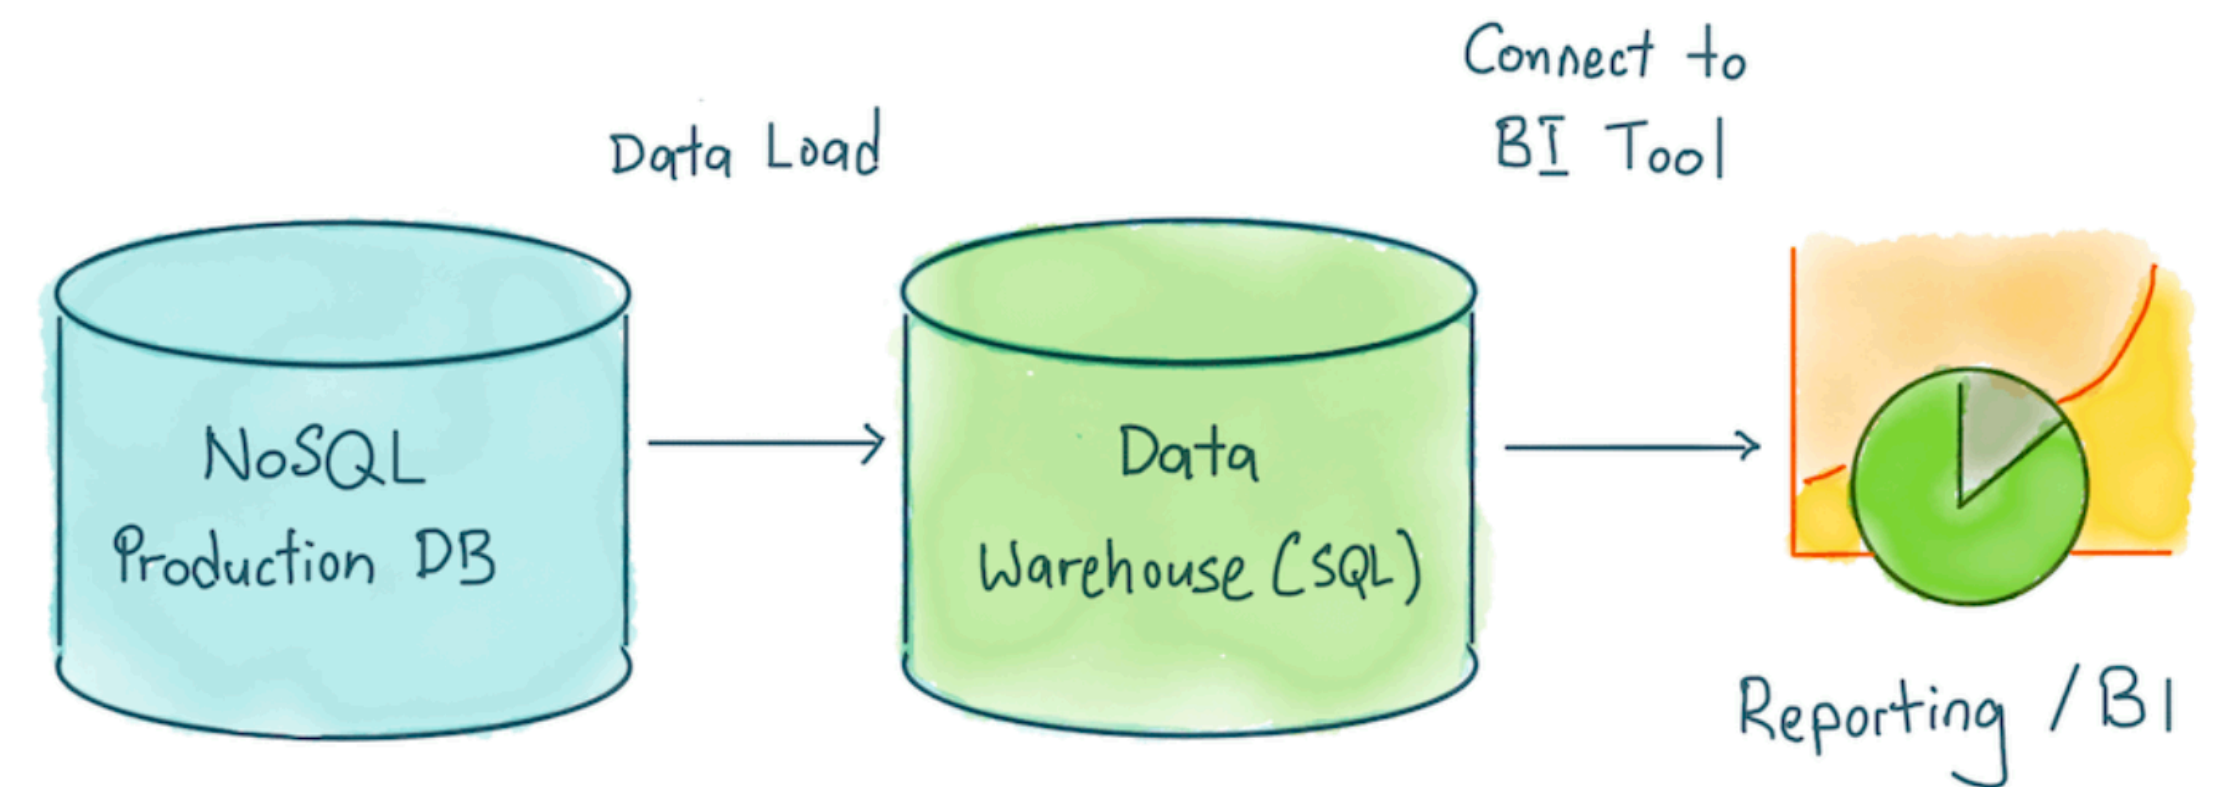


**Start by using the production database as your analytics database**

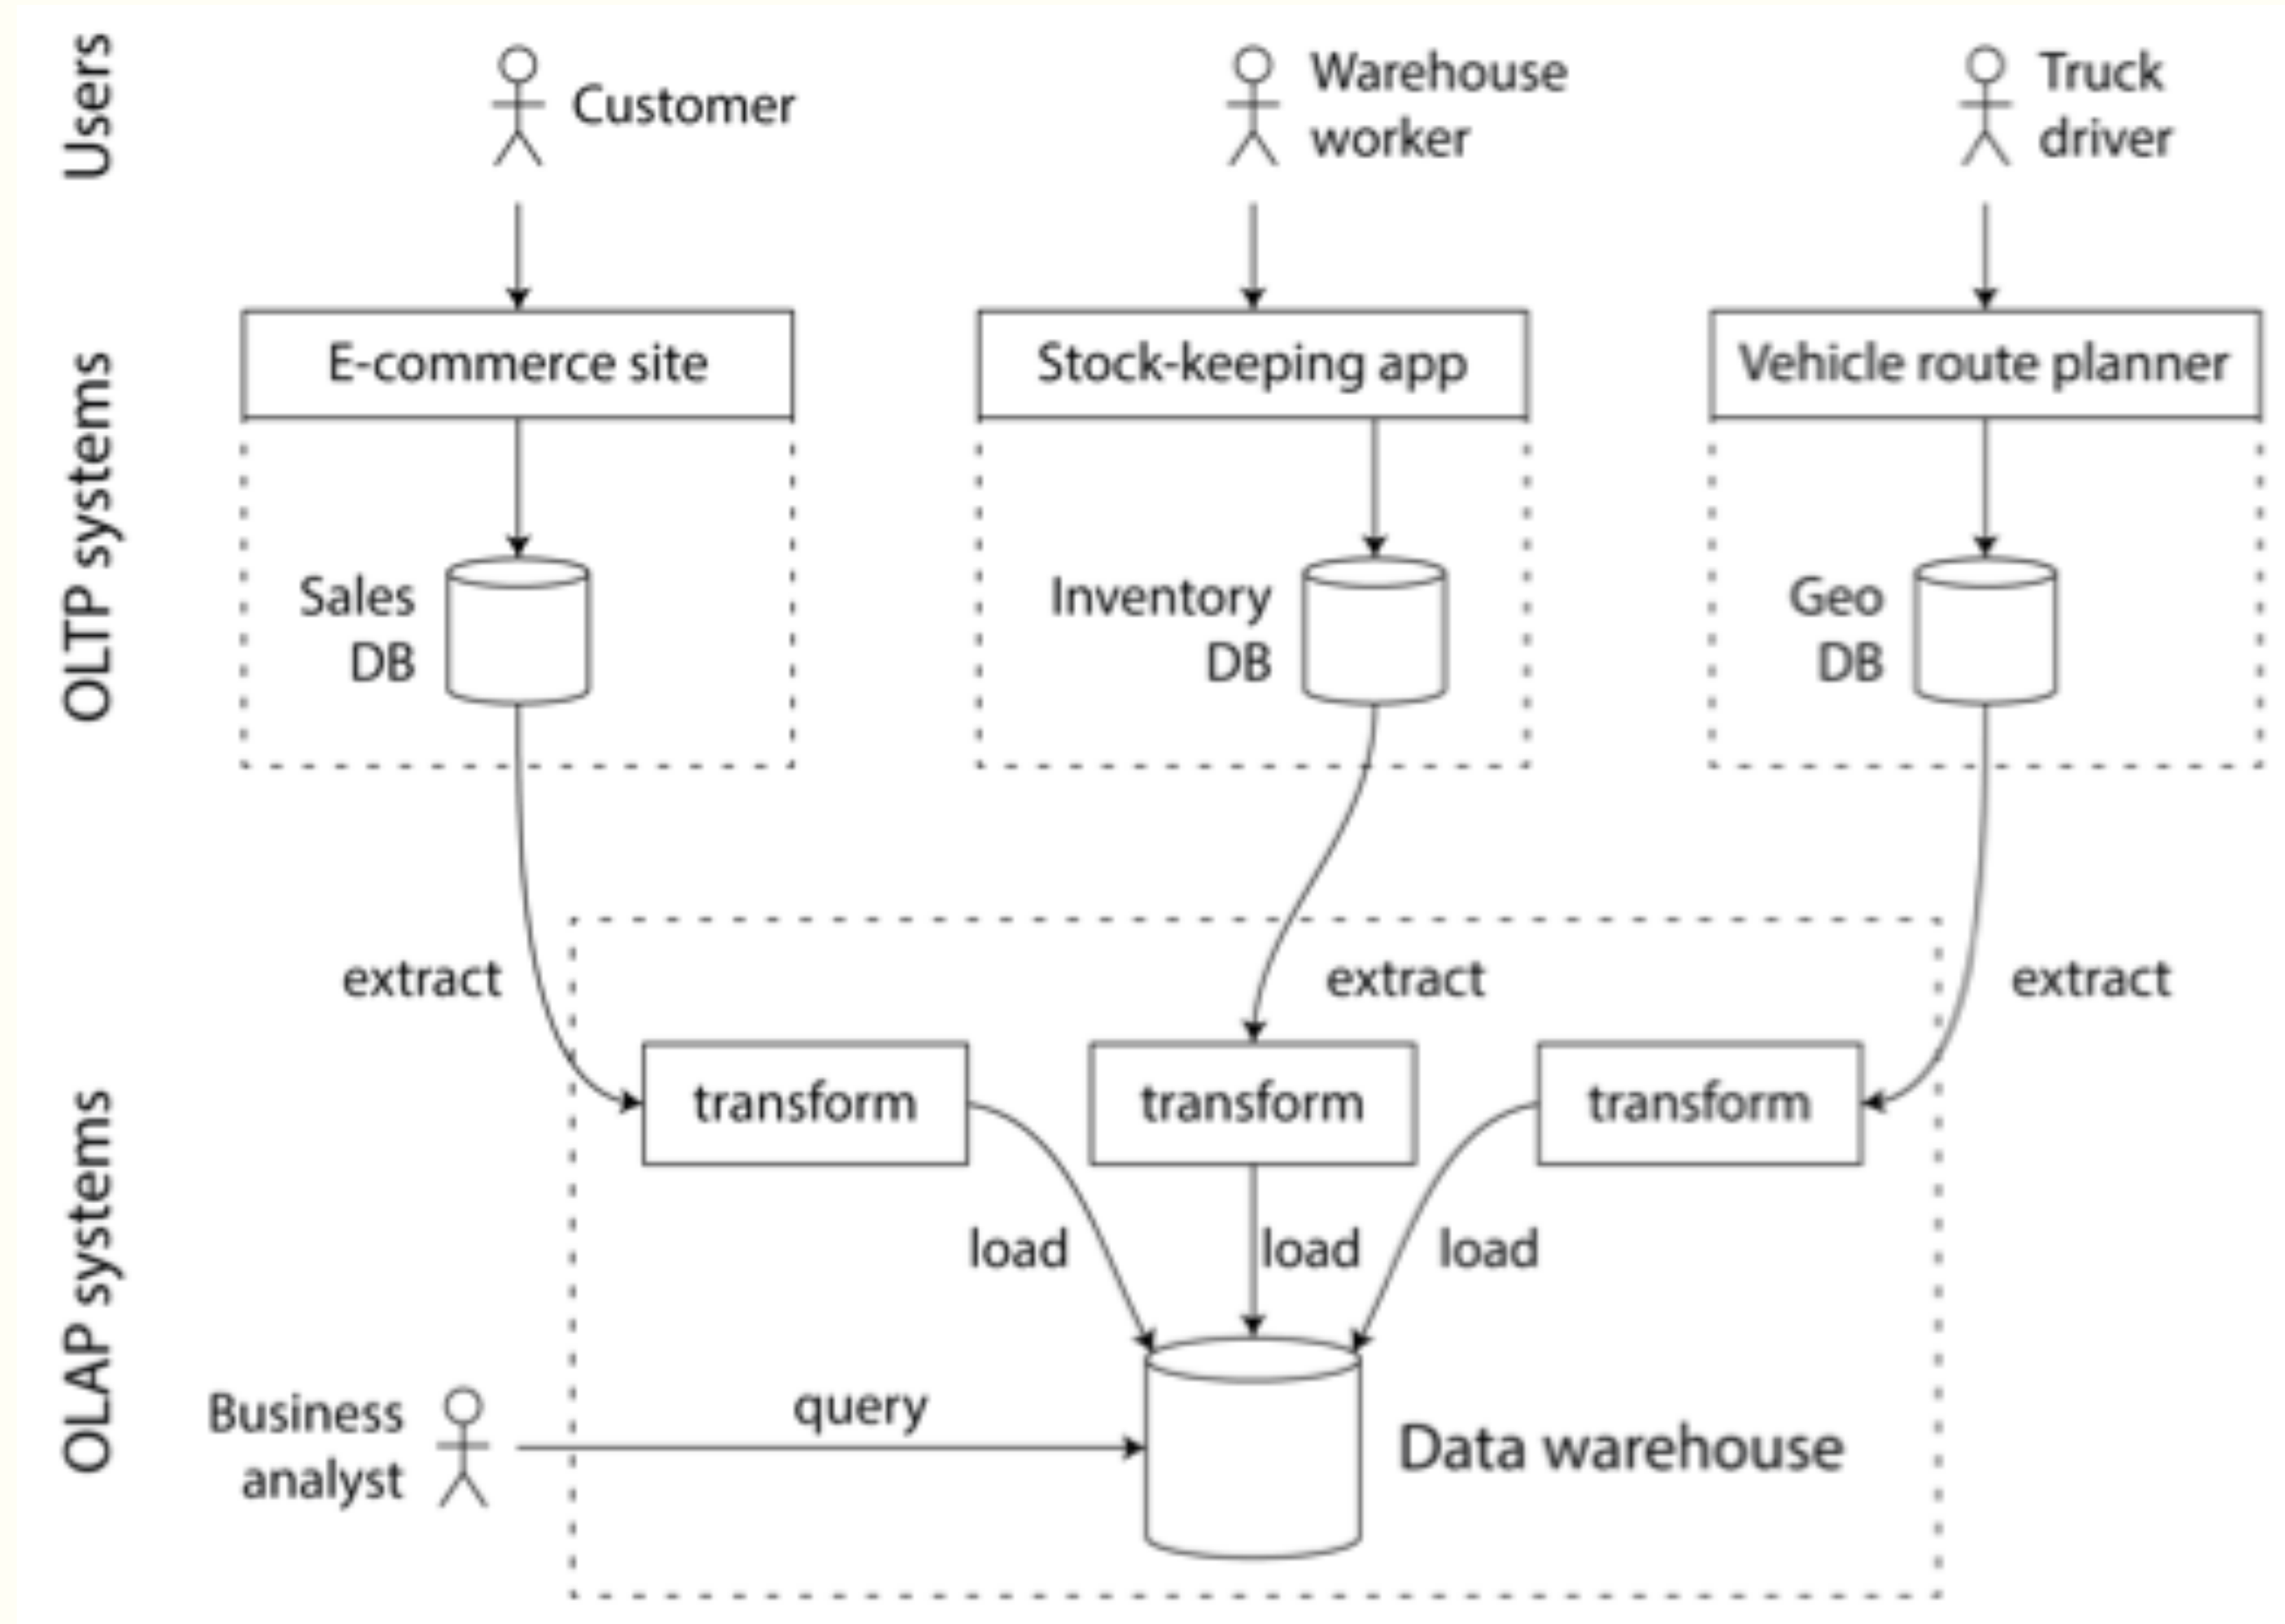


**Move to using a replica of the production database as your analytics database**

**Multiple production databases; use a warehouse instead and not affect their operation and to enable efficient analytics**



How are the  
databases  
used?





# Transactional DBs vs. Analytics DBs

## Data:

- Many single-row writes
- Current, single data

## Queries:

- Generated by user activities; 10 to 1000 users
- < 1s response time
- Short queries

## Data:

- Few large batch imports
- Years of data, many sources

## Queries:

- Generated by large reports; 1 to 10 users
- Queries run for hours
- Long, complex queries



# Columnar Databases

- Store each column separately
- Have a higher read efficiency as only a few columns of contiguous or run-encoded data need to be read
- compress better especially if the cardinality of the columns is not high thus allowing more data to be loaded into memory
- columnar data have higher sorting and indexing efficiency and may even admit multiple sort orders

# Row vs Column

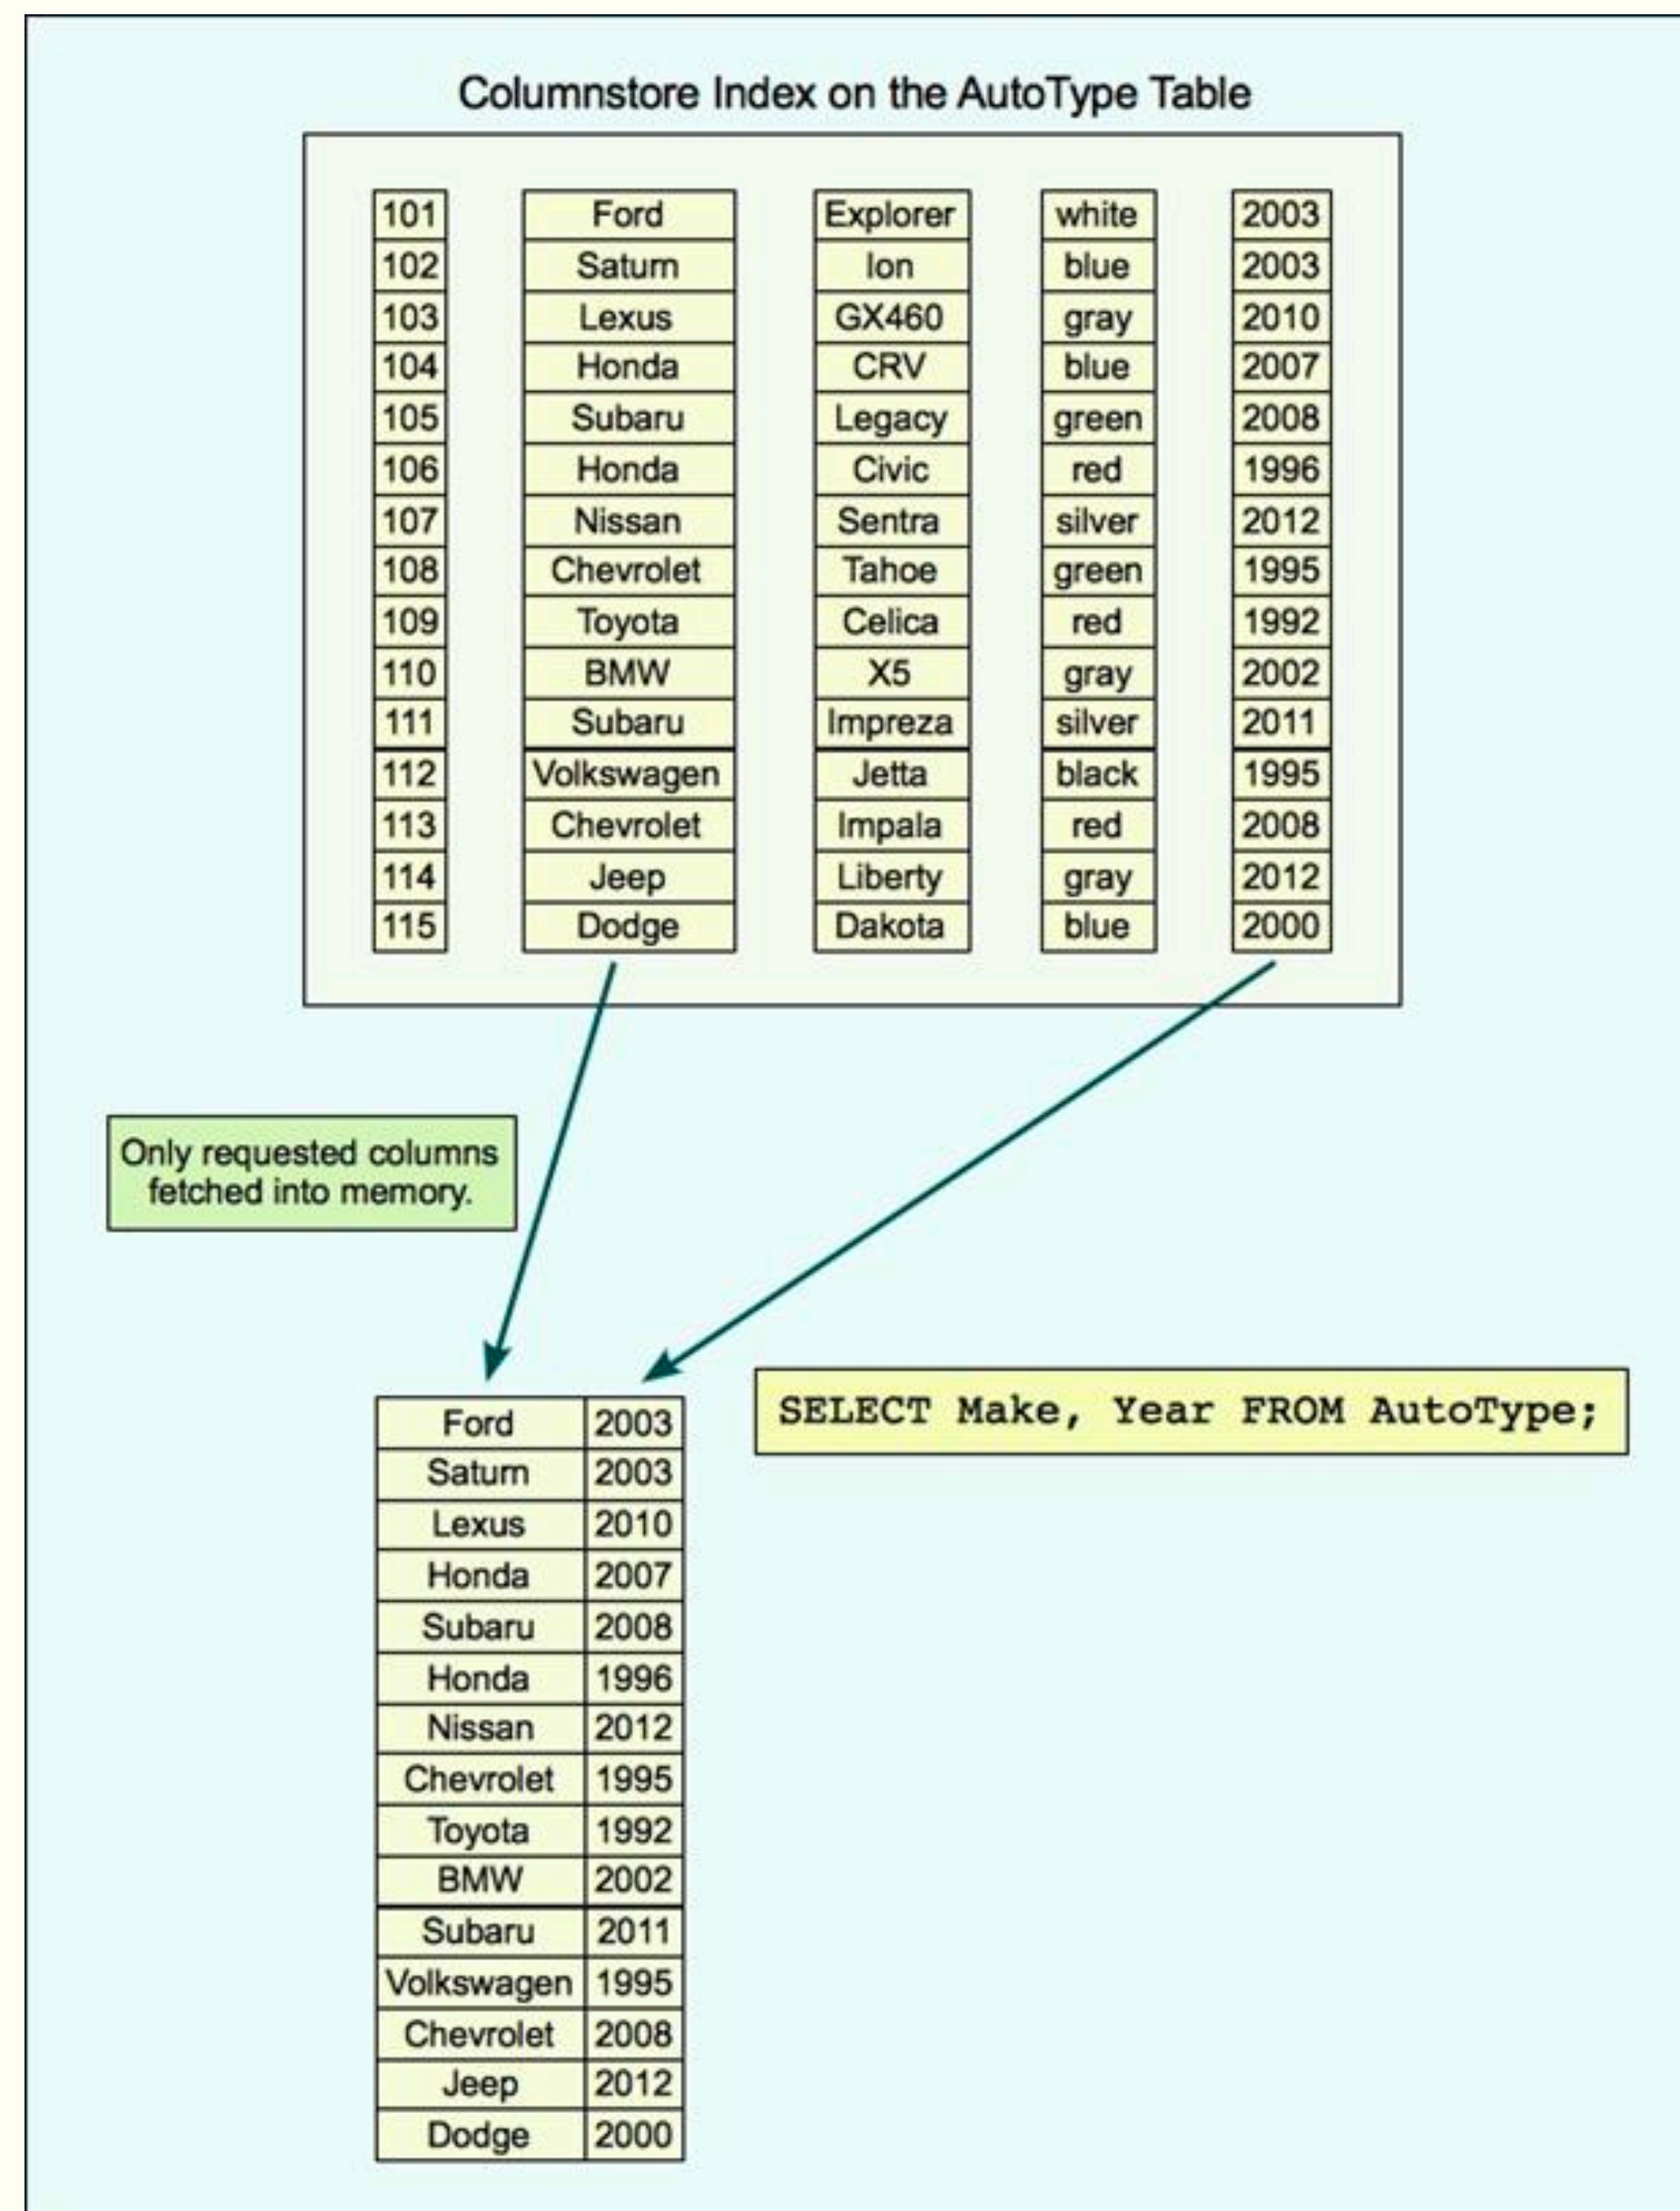
OrderId	CustomerId	ShippingCountry	OrderTotal
1	1258	US	55.25
2	5698	AUS	125.36
3	2265	US	776.95
4	8954	CA	32.16
Block 1	1, 1258, US, 55.25		
Block 2	2, 5698, AUS, 125.36		
Block 3	3, 2265, US, 776.95		
Block 4	4, 8954, CA, 32.16		

OrderId	CustomerId	ShippingCountry	OrderTotal	Customer Active
1	1258	US	55.25	TRUE
2	5698	AUS	125.36	TRUE
3	2265	US	776.95	TRUE
4	8954	CA	32.16	FALSE
Block 1	1, 2, 3, 4			
Block 2	1258, 5698, 2265, 8954			
Block 3	US, AUS, US, CA			
Block 4	55.25, 125.36, 776.95, 32.16			
Block 5	TRUE, TRUE, TRUE, FALSE			



# Column-oriented Storage

- store values from each column together in separate storage
- lends itself to compression with bitmap indexes
- compressed indexes can fit into cache and are usable by iterators
- several different sort orders can be redundantly stored
- writing is harder: updating a row touches many column files
- but you can write an in-memory front sorted store (row or column), and eventually merge onto the disk





# Bitmap Indexes

- lends itself to compression with bitmap indexes and run-length encoding. This involves choosing an appropriate sort order. The index then can be the data (great for IN and AND queries): there is no pointers to “elsewhere”
- bitwise AND/OR can be done with vector processing

Column values:  
product\_sk: 

69	69	69	69	74	31	31	31	31	29	30	30	31	31	31	68	69	69
----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----

Bitmap for each possible value:  
product\_sk = 29: 

0	0	0	0	0	0	0	0	0	1	0	0	0	0	0	0	0	0
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

  
product\_sk = 30: 

0	0	0	0	0	0	0	0	0	0	1	1	0	0	0	0	0	0
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

  
product\_sk = 31: 

0	0	0	0	0	1	1	1	1	0	0	0	1	1	1	0	0	0
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

  
product\_sk = 68: 

0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	1	0	0
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

  
product\_sk = 69: 

1	1	1	1	0	0	0	0	0	0	0	0	0	0	0	0	1	1
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---


  
product\_sk = 74: 

0	0	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

Run-length encoding:  
product\_sk = 29: 9, 1 (9 zeros, 1 one, rest zeros)  
product\_sk = 30: 10, 2 (10 zeros, 2 ones, rest zeros)  
product\_sk = 31: 5, 4, 3, 3 (5 zeros, 4 ones, 3 zeros, 3 ones, rest zeros)  
product\_sk = 68: 15, 1 (15 zeros, 1 one, rest zeros)  
product\_sk = 69: 0, 4, 12, 2 (0 zeros, 4 ones, 12 zeros, 2 ones)  
product\_sk = 74: 4, 1 (4 zeros, 1 one, rest zeros)



# ETL to ELT



**Move stuff into a warehouse or lake first!  
This is called a source refreshed table**

**Now transformations are done IN the warehouse.  
The data engineer can do the initial loading and transformations.  
Data analysts from client groups (sales, marketing)  
can do subsequent transformations**

